

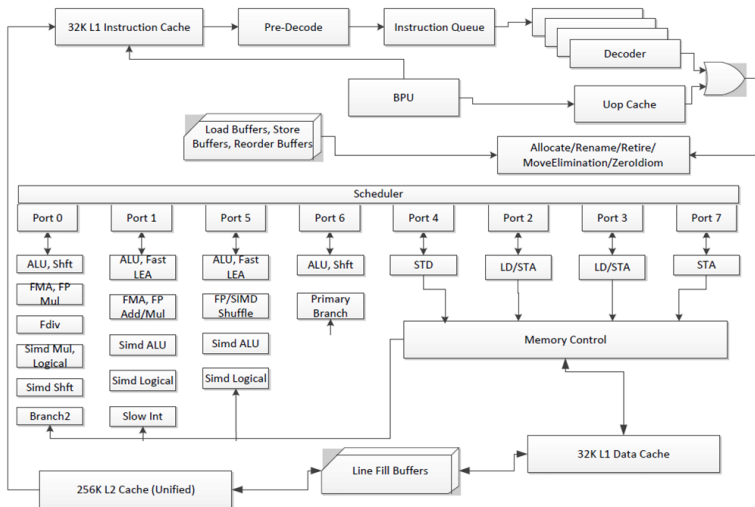
MA6252 Topics in Applied Mathematics II

Introduction to MPI


National University of Singapore
February 17th 2015

- Levels of parallelism
- MPI Overview
- Process model and language bindings
- Messages and point-to-point communication
- Non-blocking communication
- Collective communication

CPU Core Pipeline Functionality: Intel Microarchitecture Haswell



Data Parallelism - SIMD




Intel® AVX

Vector size: 256bit

Data types: 32 and 64 bit floats

VL: 4, 8, 16

Sample: X_i , Y_i 32 bit int or float



Intel® MIC

Vector size: 512bit

Data types:

32 and 64 bit integers


32 and 64bit floats

(some support for
16 bits floats)

VL: 8,16

Sample: 32 bit float

Data Parallelism - SIMD




Intel® AVX

Vector size: 256bit

Data types: 32 and 64 bit floats

VL: 4, 8, 16

Sample: X_i , Y_i 32 bit int or float



Intel® MIC

Vector size: 512bit

Data types:

32 and 64 bit integers

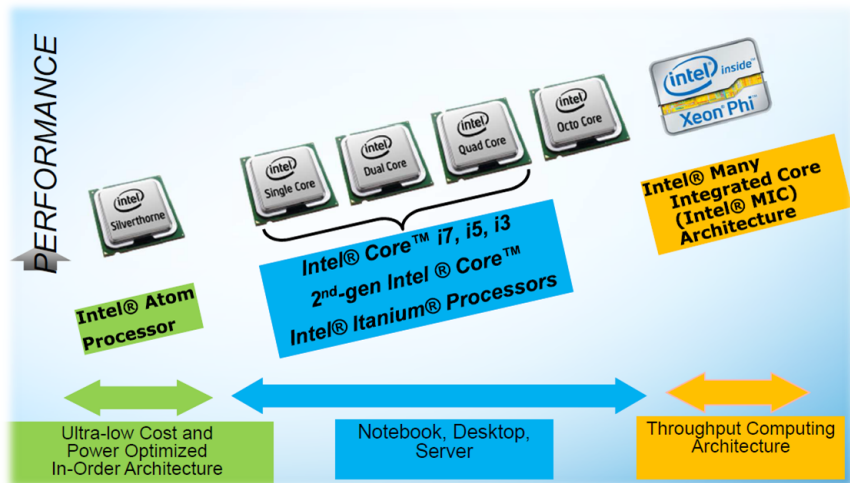
32 and 64bit floats

(some support for
16 bits floats)


VL: 8,16

Sample: 32 bit float


Shared Memory Parallelism



Shared Memory Access



conventional symmetric
multiprocessing (SMP)



non-uniform memory
access (NUMA)

Distributed Memory Machines



	JUGENE	JUQUEEN	SuperMUC
System	IBM Bluegene/P	IBM Bluegene/Q	IBM System x iDataPlex
Processor	PowerPC 450	PowerPC A2	Xeon E5-2680 8C
Clock frequency	0.85 GHz	1.6 GHz	2.8 GHz
#Nodes	73 728	24 576	9 216
#Cores	294 912	393 216	147 456
Network	3D Torus	5D Torus	Tree
GFlop/s per Watt	0.44	2.54	0.94

Is parallelism really so important?

	Peak	Factor	Paradigm
Serial	0.5 GFlops		

Is parallelism really so important?

	Peak	Factor	Paradigm
Serial	0.5 GFlops		
Processing units	1 GFlops	x2	balance between adds and mults, CPU

Is parallelism really so important?

	Peak	Factor	Paradigm
Serial	0.5 GFlops		
Processing units	1 GFlops	x2	balance between adds and mults, CPU
Instruction pipeline	6 GFlops	x6	compiler or CPU

Is parallelism really so important?

	Peak	Factor	Paradigm
Serial	0.5 GFlops		
Processing units	1 GFlops	x2	balance between adds and mults, CPU
Instruction pipeline	6 GFlops	x6	compiler or CPU
SIMD (AVX)	24 GFlops	x4	compiler, intrinsics, pragmas

Is parallelism really so important?

	Peak	Factor	Paradigm
Serial	0.5 GFlops		
Processing units	1 GFlops	x2	balance between adds and mults, CPU
Instruction pipeline	6 GFlops	x6	compiler or CPU
SIMD (AVX)	24 GFlops	x4	compiler, intrinsics, pragmas
Shared memory parallel	384 GFlops	x16	OpenMP

Is parallelism really so important?

	Peak	Factor	Paradigm
Serial	0.5 GFlops		
Processing units	1 GFlops	x2	balance between adds and mults, CPU
Instruction pipeline	6 GFlops	x6	compiler or CPU
SIMD (AVX)	24 GFlops	x4	compiler, intrinsics, pragmas
Shared memory parallel	384 GFlops	x16	OpenMP
(Accelerators, e.g. GPU)			CUDA, OpenCL


Is parallelism really so important?

	Peak	Factor	Paradigm
Serial	0.5 GFlops		
Processing units	1 GFlops	x2	balance between adds and mults, CPU
Instruction pipeline	6 GFlops	x6	compiler or CPU
SIMD (AVX)	24 GFlops	x4	compiler, intrinsics, pragmas
Shared memory parallel	384 GFlops	x16	OpenMP
(Accelerators, e.g. GPU)			CUDA, OpenCL
Distributed memory parallel	3.84 PFlops	x10 000	MPI

- MPI Overview
 - one program on several processors
 - work and data distribution
 - the communication network
- Process model and language bindings
- Messages and point-to-point communication
- Non-blocking communication
- Virtual topologies
- Collective communication


The Message-Passing Programming Paradigm

Sequential Programming Paradigm




The Message-Passing Programming Paradigm

Sequential Programming Paradigm




Message-Passing Programming Paradigm



The Message-Passing Programming Paradigm


Each processor in a message passing program runs a **sub-program**

- written in a conventional sequential language, e.g., C(++) or Fortran
- typically the same on each processor (SPMD)
- the variables of each sub-program have
 - the same name
 - but different locations (distributed memory) and different data!
 - i.e., all variables are private
- communicate via special send & receive routines (**message passing**)




Data and Work Distribution

- the value of **myrank** is returned by special library function
- the system of **size** processes is started by special MPI initialization program (mpirun or mpiexec)
- all distribution decisions are based on **myrank**
- i.e., which process works on which data



Messages



- Messages are packets of data moving between sub-programs
- Necessary information for the message passing system:
 - sending process
 - source location
 - source data type
 - source data size
 - receiving process
 - destination location
 - destination data type
 - destination buffer size

- A sub-program needs to be connected to a message passing system
- A message passing system is similar to:
 - mail box
 - phone line
 - fax machine
 - etc.
- MPI:
 - sub-program must be linked with an MPI library
 - the total program (i.e., all sub-programs of the program) must be started with the MPI startup tool

- Messages need to have addresses to be sent to
- Addresses are similar to:
 - mail addresses
 - phone number
 - fax number
 - etc.
- MPI: addresses are ranks of the MPI processes (sub-programs)

Point-to-Point communication

- Simplest form of message passing
- One process send a message to another
- Different types of point-to-point communication:
 - non-buffered = synchronous send
 - buffered = asynchronous send

Synchronous/Asynchronous Sends

Synchronous Sends

- The sender gets an information that the message is received
- Analog to the beep or okay-sheet of a fax

Synchronous/Asynchronous Sends

Synchronous Sends

- The sender gets an information that the message is received
- Analog to the beep or okay-sheet of a fax

Buffered = Asynchronous Sends

- We only know when the message has left

Synchronous/Asynchronous Sends – Examples

Process 0

```
send    ( &a, 1, 1 );  
receive( &b, 1, 1 );
```

Process 1

```
send    ( &a, 1, 0 );  
receive( &b, 1, 0 );
```

Synchronous/Asynchronous Sends – Examples

Process 0

```
send    ( &a, 1, 1 );  
receive( &b, 1, 1 );
```

Process 1

```
send    ( &a, 1, 0 );  
receive( &b, 1, 0 );
```

Depending on the implementation of send and receive this might lead to a **deadlock!**

Synchronous/Asynchronous Sends – Examples

Process 0

```
send    ( &a, 1, 1 );  
receive( &b, 1, 1 );
```

Process 1

```
send    ( &a, 1, 0 );  
receive( &b, 1, 0 );
```

Depending on the implementation of send and receive this might lead to a **deadlock!**

Process 0

```
receive( &b, 1, 1 );  
send    ( &a, 1, 1 );
```

Process 1

```
receive( &b, 1, 0 );  
send    ( &a, 1, 0 );
```

Synchronous/Asynchronous Sends – Examples

Process 0

```
send    ( &a, 1, 1 );  
receive( &b, 1, 1 );
```

Process 1

```
send    ( &a, 1, 0 );  
receive( &b, 1, 0 );
```

Depending on the implementation of send and receive this might lead to a **deadlock!**

Process 0

```
receive( &b, 1, 1 );  
send    ( &a, 1, 1 );
```

Process 1

```
receive( &b, 1, 0 );  
send    ( &a, 1, 0 );
```

Even with buffered send this code will lead to a **deadlock!**

Blocking Operations

- Operations are local activities, e.g.,
 - sending (a message)
 - receiving (a message)
- Some operations may **block** until another process acts:
 - synchronous send operation **blocks until** receive is posted
 - receive operation **blocks until** message is sent
- Relates to the completion of an operation
- Blocking subroutine returns only when the operation has completed

Non-Blocking Operations

- Non-blocking operation: returns immediately and allow the sub-program to perform other work.
- At some later time the sub-program must **test** or **wait** for the completion of the non-blocking operation.
- All non-blocking operations must have matching wait (or test) operations. (Some system or application resources can be freed only when the non-blocking operation is completed.)
- A non-blocking operation immediately followed by a matching wait is equivalent to a blocking operation
- Non-blocking operations are not the same as sequential subroutine calls: the operation may continue while the application executes the next statements!

Collective Communications

- Collective communication routines are higher level routines.
- Several processes are involved at a time
- May allow optimized internal implementations, e.g., tree based algorithms
- Can be built out of point-to-point communications

Broadcast

- A one-to-many communication

Collective Communications (cont'd)

Broadcast

- A one-to-many communication

Reduction Operations

- Combine data from several processes to produce a single result

Collective Communications (cont'd)

Broadcast

- A one-to-many communication

Reduction Operations

- Combine data from several processes to produce a single result

Barriers

- Synchronize processes

- MPI Overview
- Process model and language bindings
 - starting several MPI processes
- Messages and point-to-point communication
- Non-blocking communication
- Derived data types
- Virtual topologies
- Collective communication

Compilation

```
g++ ... -L<mpi_library_path> -l<mpi_library_name>  
      -I<mpi_include_path> ...  
mpiCC ...
```

Compilation

```
g++ ... -L<mpi_library_path> -l<mpi_library_name>  
      -I<mpi_include_path> ...  
mpiCC ...
```

Header files

```
#include <mpi.h>
```

Compilation/Header Files/Function Format

Compilation

```
g++ ... -L<mpi_library_path> -l<mpi_library_name>  
      -I<mpi_include_path> ...  
mpiCC ...
```

Header files

```
#include <mpi.h>
```

MPI Function Format

```
error = MPI_Xxxxx( parameter, ... );  
MPI_Xxxxxx( parameter, ... );
```


MPI Function Format Details

- MPI_.... namespace is reserved for MPI constants and routines, i.e., application routines and variable names must not begin with MPI_.
- Output arguments in C(++) are handled via passing-by-pointer:

Function definitions

```
MPI_Comm_rank( ..., int *rank );  
MPI_Recv( ..., MPI_Status *status );
```

Usage in your code

```
main( ... ) {  
    int myrank;  
    MPI_Status recv_status;  
    MPI_Comm_rank( ..., &myrank );  
    MPI_Recv( ..., &recv_status );  
}
```

Initializing MPI

- Should be the very first function call inside the main function
- Must be the very first MPI function call

Function definition


```
MPI_Init( int *argc, char ***argv );
```

Usage in your code

```
#include <mpi.h>
int main( int argc, char** argv )
{
    MPI_Init( &argc, &argv );
    ...
}
```

Starting the MPI Program


- Start mechanism is implementation dependent
- `mpirun -np number_of_processes ./executable` (most implementations)
- `mpiexec -n number_of_processes ./executable` (with MPI-2 standard)



- The parallel MPI processes exist at least after `MPI_Init` was called.

Communicator MPI_COMM_WORLD


- All processes (= sub-programs) of one MPI program are combined in the communicator MPI_COMM_WORLD
- MPI_COMM_WORLD is a predefined **handle** in `mpi.h`
- Each process has its own **rank** in a communicator:
 - starting with 0
 - ending with size-1



Rank

- The rank identifies different processes
- The rank is the basis for any work and data distribution

```
int MPI_Comm_rank( MPI_Comm comm, int *rank );
```



- How many processes are contained within a communicator?

```
int MPI_Comm_size( MPI_Comm comm, int *size );
```

```
int MPI_Finalize();
```

- **MUST** be called last by all processes
- After `MPI_Finalize`:
 - Further MPI-calls are forbidden
 - Especially re-initialization with `MPI_Init` is forbidden

Example: MPI Hello World (I)

```
#include <mpi.h>

int main( int argc, char **argv )
{
    // Definition of the variables
    int size; //The total number of processes
    int rank; //The rank/number of this process

    // MPI initialization
    MPI_Init( &argc, &argv );

    // Determining the number of CPUs and the rank for each CPU
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );

    ...
}
```


Example: MPI Hello World (II)

```
...

// 'Hello World' output for CPU 0
if( rank == 0 )
    cout << " Hello World" << endl;

// Output of the own rank and size of each CPU
cout << " I am CPU " << rank << " of " << size << " CPUs" << endl;

// MPI finalizations
MPI_Finalize();

return 0;
}
```

Example: MPI Hello World (III)

Possible output of the program:


```
I am CPU 2 of 4 CPUs  
Hello World  
I am CPU 0 of 4 CPUs  
I am CPU 3 of 4 CPUs  
I am CPU 1 of 4 CPUs
```

Note: The output of the program is non-deterministic!

- MPI Overview
- Process model and language bindings
- Messages and point-to-point communication
 - the MPI processes can communicate
- Non-blocking communication
- Derived data types
- Virtual topologies
- Collective communication

Messages

- A message contains a number of elements of some particular datatype
- MPI datatypes:
 - Basic datatype
 - Derived datatypes



- Derived datatypes can be built up from basic or derived data types
- Data type handles are used to describe the type of the data in the memory.
- Example: message with 5 integers


2345	654	96574	-12	7676
------	-----	-------	-----	------

MPI Basic Datatypes

MPI_CHAR	char	Treated as printable character
MPI_SHORT	signed short int	
MPI_INT	signed int	
MPI_LONG	signed long int	
MPI_LONG_LONG	signed long long	
MPI_SIGNED_CHAR	signed char	Treated as integral value
MPI_UNSIGNED_CHAR	unsigned char	Treated as integral value
MPI_UNSIGNED_SHORT	unsigned short int	
MPI_UNSIGNED	unsigned int	
MPI_UNSIGNED_LONG	unsigned long int	
MPI_UNSIGNED_LONG_LONG	unsigned long long	
MPI_FLOAT	float	
MPI_DOUBLE	double	
MPI_LONG_DOUBLE	long double	
MPI_BYTE		
MPI_PACKED		

Point-to-Point Communication

- Communication between two processes
- Source process sends message to destination process
- Communication takes place within a communication, e.g., `MPI_COMM_WORLD`
- Processes are identified by their ranks in the communicator



Sending a Message

```
int MPI_Send( void *buf, int count, MPI_Datatype datatype,  
              int dest, int tag, MPI_Comm comm );
```

- `buf` is the starting point of the message with `count` elements
- `dest` is the rank of the destination process within the communicator `comm`
- `tag` is an additional nonnegative integer piggyback information, additionally transferred with the message
- The tag can be used by the program to distinguish different types of messages

Receiving a Message

```
int MPI_Recv( void *buf, int count, MPI_Datatype datatype,  
             int source, int tag, MPI_Comm comm,  
             MPI_Status *status );
```

- buf, count, datatype describe the receive buffer
- Receiving the message sent by process with rank source in comm
- Envelope information is returned in status
- Only messages with matching tag are received

Requirements for Point-to-Point Communications

For a communication to succeed:


- Sender must specify a valid destination rank
- Receiver must specify a valid source rank
- The communicator must be the same
- Tags must match
- Message datatypes must match
- Receiver's buffer must be large enough

- Receiver can wildcard
- To receive from any source \Rightarrow `source = MPI_ANY_SOURCE`
- To receive from any tag \Rightarrow `tag = MPI_ANY_TAG`
- Actual source and tag are returned in the receiver's status parameter

Communication Envelope

Envelope information is returned from
MPI_Recv in status

```
status.MPI_SOURCE  
status.MPI_TAG  
status.MPI_ERROR  
count via MPI_Get_count()
```



```
int MPI_Get_count( MPI_Status *status,  
                  MPI_Datatype datatype, int *count );
```

Communication Modes

Send communication modes	⇒	
- synchronous send	⇒	MPI_Ssend
- buffered [asynchronous] send	⇒	MPI_Bsend
- standard send	⇒	MPI_Send
- ready send	⇒	MPI_Rsend

Receiving all modes	⇒	MPI_Recv
---------------------	---	----------

Communication Modes – Definitions


Sender mode	Definition	Notes
Synchronous send MPI_Ssend	Only completes when the receive has started	
Buffered send MPI_Bsend	Always completes (unless an error occurs), irrespective of receiver	needs application-defined buffer to be declared with <code>MPI_Buffer_attach</code>
Standard send MPI_Send	Either synchronous or buffered	Uses an internal buffer
Ready send MPI_Rsend	May be started only if the matching receive is already posted!	Highly dangerous!
Receive MPI_Recv	Completes when a message has arrived	same routine for all communication modes

Rules for the Communication Modes

- Standard send (**MPI_Send**)
 - minimal transfer time
 - may block due to synchronous mode
 - → risks with synchronous send
- Synchronous send (**MPI_Ssend**)
 - risk of deadlock
 - risk of serialization
 - risk of waiting → idle time
 - high latency / best bandwidth
- Buffered send (**MPI_Bsend**)
 - low latency / bad bandwidth
- Ready send (**MPI_Rsend**)
 - use **never**, except you have a 200% guarantee that Recv is already called in the current version and all future versions of your code

Message Order Preservation

- Rule for messages on the same connection, i.e., same communicator, source, and destination rank:
- **Messages do not overtake each other**
- This is true even for non-synchronous sends.




- If both receives match both messages, the order is preserved.

- MPI Overview
- Process model and language bindings
- Messages and point-to-point communication
- Non-blocking communication
 - - to avoid idle time and deadlocks
- Derived data types
- Virtual topologies
- Collective communication

Deadlock

- Code in each MPI process:

```
// This will block and never return, because MPI_Recv  
// cannot be called in the right-hand MPI process  
MPI_Ssend( ..., right_rank, ... );  
MPI_Recv( ..., left_rank, ... );
```



- Same problem with standard send mode (MPI_Send) if MPI implementation chooses synchronous protocol

Non-Blocking Communications

Separate communication into three phases:

- Initiate non-blocking communication
 - returns immediately
 - routine name starting with `MPI_I...`
- Do some work (perhaps involving other communications?)
- Wait for non-blocking communication to complete

Non-Blocking Examples

Non-blocking send

MPI_Isend(...)
doing some other work
MPI_Wait(...)



Non-Blocking Examples

Non-blocking send

MPI_Isend(...)
doing some other work
MPI_Wait(...)




Non-blocking receive

MPI_Irecv(...)
doing some other work
MPI_Wait(...)



Non-Blocking Receive

- Initiate non-blocking receive
In the ring example: Initiate non-blocking receive from left neighbor
- Do some work:
in the ring example: Sending the message to the right neighbor
- Now, the message transfer can be completed
- Wait for non-blocking receive to complete



Request handles

- are used for non-blocking communication
- must be stored in local variables (MPI_Request)
- the value
 - **is generated** by a non-blocking communication routine
 - **is used** (and freed) in the MPI_Wait routine

Non-Blocking Synchronous Send

```
int MPI_Issend( void* buf, int count,  
               MPI_Datatype datatype, int dest,  
               int tag, MPI_Comm comm,  
               MPI_Request *request );  
  
int MPI_Wait( MPI_Request *request, MPI_Status *status );
```

- buf must not be accessed between Issend and Wait
(In MPI-2.2, this restriction is relaxed to “must not be modified”)
- “Issend + Wait directly after Issend” is equivalent to blocking call (Ssend)
- status is not used in Issend, but in Wait (with send: nothing returned)

Non-Blocking Receive

```
int MPI_Irecv( void* buf, int count,  
              MPI_Datatype datatype, int source,  
              int tag, MPI_Comm comm,  
              MPI_Request *request);
```

```
int MPI_Wait( MPI_Request *request, MPI_Status *status );
```

- buf must not be used between Irecv and Wait

Blocking and Non-Blocking

- Send and receive can be blocking or non-blocking
- A blocking send can be used with a non-blocking receive, and vice-versa.
- Non-blocking sends can use any mode
 - standard – `MPI_Isend`
 - synchronous – `MPI_Issend`
 - buffered – `MPI_Ibsend`
 - ready – `MPI_Irsend`
- Synchronous mode affects completion, i.e., `MPI_Wait`/`MPI_Test`, not initiation, i.e., `MPI_I...`
- The non-blocking operation immediately followed by a matching wait is equivalent to the blocking operation

```
int MPI_Wait( MPI_Request *request, MPI_Status *status );  
int MPI_Test( MPI_Request *request, int *flag,  
              MPI_Status *status );
```


One must

- Wait or
- loop with Test until request is completed, i.e., `flag == 1`

Multiple Non-Blocking Communications

You have several request handles:

- Wait or test for completion of **one** message
MPI_Waitany / MPI_Testany
- Wait or test for completion of **all** messages
MPI_Waitall / MPI_Testall *)
- Wait or test for completion of **as many** messages as possible
MPI_Waitsome / MPI_Testsome *)



*) Each status contains an additional error field. This field is only used if MPI_ERR_IN_STATUS is returned (also valid for send operations)

- MPI Overview
- Process model and language bindings
- Messages and point-to-point communication
- Non-blocking communication
- Derived data types
- Virtual topologies
- Collective communication
 - - e.g., broadcast

Collective communication

- Communications involving a group of processes
- Called by all processes in a communicator
- Examples:
 - Barrier synchronization
 - Broadcast, scatter, gather
 - Global sum, global maximum, etc.

Characteristics of Collective communication

- Collective action over a communicator
- All processes of the communicator must communicate, i.e., must call the collective routine
- Synchronization may or may not occur, therefore all processes must be able to start the collective routine
- All collective operations are blocking
- No tags
- Receive buffers must have exactly the same size as send buffers

Barrier Synchronization

```
MPI_Barrier( MPI_Comm comm );
```


MPI_Barrier is normally never needed:

- all synchronization is done automatically by the data communication
 - a process cannot continue before it has the data that it needs
- if used for debugging
 - please guarantee, that it is removed in production
- for profiling: to separate time measurement of
 - Load imbalance of computation [MPI_Wtime(); MPI_Barrier(); MPI_Wtime();]
 - communication epochs [MPI_Wtime(); MPI_Allreduce(); ...; MPI_Wtime();]
- if used for synchronizing external communication (e.g., I/O)
 - exchanging tokens may be more efficient and scalable than a barrier on MPI_COMM_WORLD

Broadcast


```
int MPI_Bcast( void* buffer, int count,  
               MPI_Datatype datatype,  
               int root, MPI_Comm comm );
```

e.g. root=2




Scatter

```
int MPI_Scatter( void* sendbuf, int sendcount,  
                MPI_Datatype sendtype, void* recvbuf,  
                int recvcount, MPI_Datatype recvtype,  
                int root, MPI_Comm comm );
```



Gather

```
int MPI_Gather( void* sendbuf, int sendcount,  
               MPI_Datatype sendtype, void* recvbuf,  
               int recvcount, MPI_Datatype recvttype,  
               int root, MPI_Comm comm );
```



Global Reduction Operations


- To perform a global reduce operation across all members of a group
- $d_0 \circ d_1 \circ d_2 \circ \dots \circ d_{s-2} \circ d_{s-1}$
 - d_i = data in process rank i
 - single variable, or
 - vector
 - \circ = associative operation
 - Example:
 - global sum or product
 - global maximum or minimum
 - global user-defined operation
 - floating point rounding may depend on usage of associative law
 - $[(d_0 \circ d_1) \circ (d_2 \circ d_3)] \circ [\dots \circ (d_{s-2} \circ d_{s-1})]$
 - $(((((d_0 \circ d_1) \circ d_2) \circ d_3) \circ \dots) \circ d_{s-2}) \circ d_{s-1})$

Predefined Reduction Operation Handles

Predefined operation handle	Function
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI LAND	Logical AND
MPI_BAND	Bitwise AND
MPI_LOR	Logical OR
MPI BOR	Bitwise OR
MPI_LXOR	Logical exclusive OR
MPI_BXOR	Bitwise exclusive OR
MPI_MAXLOC	Maximum and location of the maximum
MPI_MINLOC	Minimum and location of the minimum

MPI_Reduce

```
int MPI_Reduce( void* sendbuf, void* recvbuf,  
               int count, MPI_Datatype datatype,  
               MPI_Op op, int root, MPI_Comm comm );
```



Example of Global Reduction

- Global integer sum
- Sum of all inbuf values should be returned in resultbuf
- ```
root = 0;
MPI_reduce(&inbuf, &resultbuf, 1, MPI_INT, MPI_SUM,
 root, MPI_COMM_WORLD);
```
- The result is only placed in resultbuf at the root process


# Variants of Reduction Operations

- `MPI_Allreduce`
  - no root
  - returns the result in all processes
- `MPI_Reduce_scatter`
  - result vector of the reduction operation is scattered to the processes into the real result buffers
- `MPI_Scan`
  - prefix reduction
  - result at process with rank  $i$  := reduction of inbuf-values from rank 0 to rank  $i$




# MPI\_Allreduce

```
int MPI_Allreduce(void* sendbuf, void* recvbuf,
 int count, MPI_Datatype datatype,
 MPI_Op op, MPI_Comm comm);
```



# MPI\_Scan

```
int MPI_Scan(void* sendbuf, void* recvbuf,
 int count, MPI_Datatype datatype,
 MPI_Op op, MPI_Comm comm);
```



- MPI Overview
- Process model and language bindings
- Messages and point-to-point communication
- Non-blocking communication
- Derived data types
- Virtual topologies
- Collective communication
- **Serialization**
  - - how to order the execution of MPI processes

## Question

How can MPI processes be serialized?

# Solution 1

```
if(rank == 0) /* ... */
MPI_Barrier(MPI_COMM_WORLD);
if(rank == 1) /* ... */
MPI_Barrier(MPI_COMM_WORLD);
// ...
```

# Solution 1

```
if(rank == 0) /* ... */
MPI_Barrier(MPI_COMM_WORLD);
if(rank == 1) /* ... */
MPI_Barrier(MPI_COMM_WORLD);
// ...
```

```
MPI_Comm_size(MPI_COMM_WORLD, &size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);

for(int i=0; i<size; ++i) {
 if(i == rank) /* ... */
 MPI_Barrier(MPI_COMM_WORLD);
}
```

## Solution 2

- Each process waits for a message from the processes with  $\text{rank}-1$  (blocking receive)
- Do some work
- Send a message to the process with rank  $\text{rank}+1$

- MPI Overview
- Process model and language bindings
- Messages and point-to-point communication
- Non-blocking communication
- Derived data types
- Virtual topologies
  - - a multi-dimensional process naming scheme
- Collective communication

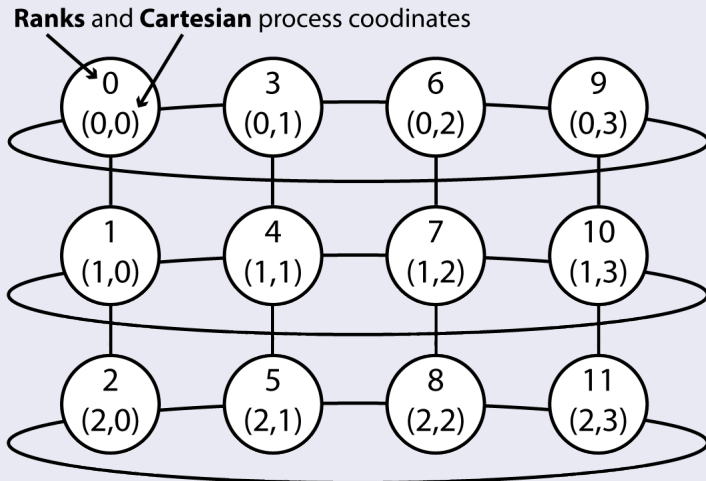


- Convenient process naming
- Naming scheme to fit the communication pattern
- Simplifies writing of code
- Can allow MPI to optimize communications

# How to use a Virtual Topology

- Creating a topology produces a new communicator
- MPI provides mapping functions
  - to compute process ranks, based on the topology naming scheme
  - and vice versa.

## Example – A 2-dimensional Cylinder

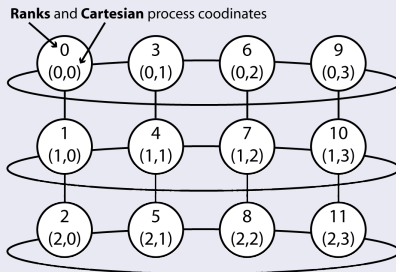


- Cartesian Topologies
  - each process is connected to its neighbor in a virtual grid
  - boundaries can be cyclic, or not
  - processes are identified by Cartesian coordinates
  - of course, communication between any two processes is still allowed
- Graph Topologies
  - general graphs
  - not covered here


# Creating a Cartesian Virtual Topology

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims,
 int *dims, int *periods, int reorder,
 MPI_Comm *comm_cart);
```

```
comm_old = MPI_COMM_WORLD
ndims = 2
dim = { 3, 4 }
periods = { 0, 1 }
reorder = see next slide
```




## Example – A 2-dimensional Cylinder



- Ranks in `comm` and `comm_cart` may differ, if `reorder == 1`
- This reordering can allow MPI to optimize communications

# Cartesian Mapping Functions


Mapping ranks to process grid coordinates



```
MPI_Cart_coords(MPI_Comm comm, int rank,
 int maxdims, int *coords);
```

# Cartesian Mapping Functions

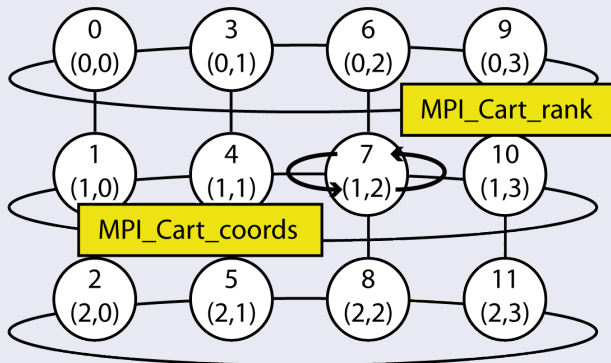
Mapping process grid coordinates to ranks



```
MPI_Cart_rank(MPI_Comm comm, int *coords, int *rank);
```



# Own Coordinates



Each process gets its own coordinates with

```
MPI_Comm_rank (comm_cart, &my_rank);
MPI_Cart_coords(comm_cart, my_rank, maxdims, my_coords);
```


# Cartesian Mapping Functions

- Computing ranks of neighboring processes

```
MPI_Cart_shift(MPI_Comm comm, int direction, int disp,
 int *rank_source, int *rank_dest);
```

- Returns MPI\_PROC\_NULL if there is no neighbor
- MPI\_PROC\_NULL can be used as source or destination rank in each communication → Then, this communication will be a noop!

# MPI\_Cart\_shift – Example



`MPI_Cart_shift( cart, direction, displace, rank_source, rank_dest );`  
example of  
process rank=7

|   |    |   |    |
|---|----|---|----|
| 0 | +1 | 6 | 8  |
| 1 | +1 | 4 | 10 |